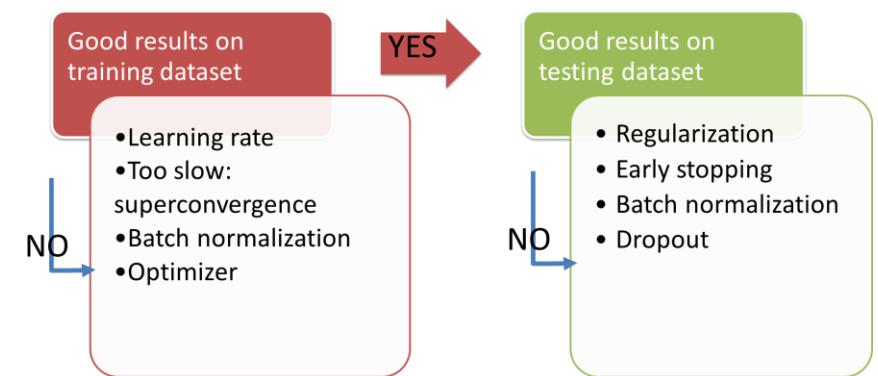


# Lecture 4-3 Optimizer: Gradient Descent Optimization

Tian Sheuan Chang

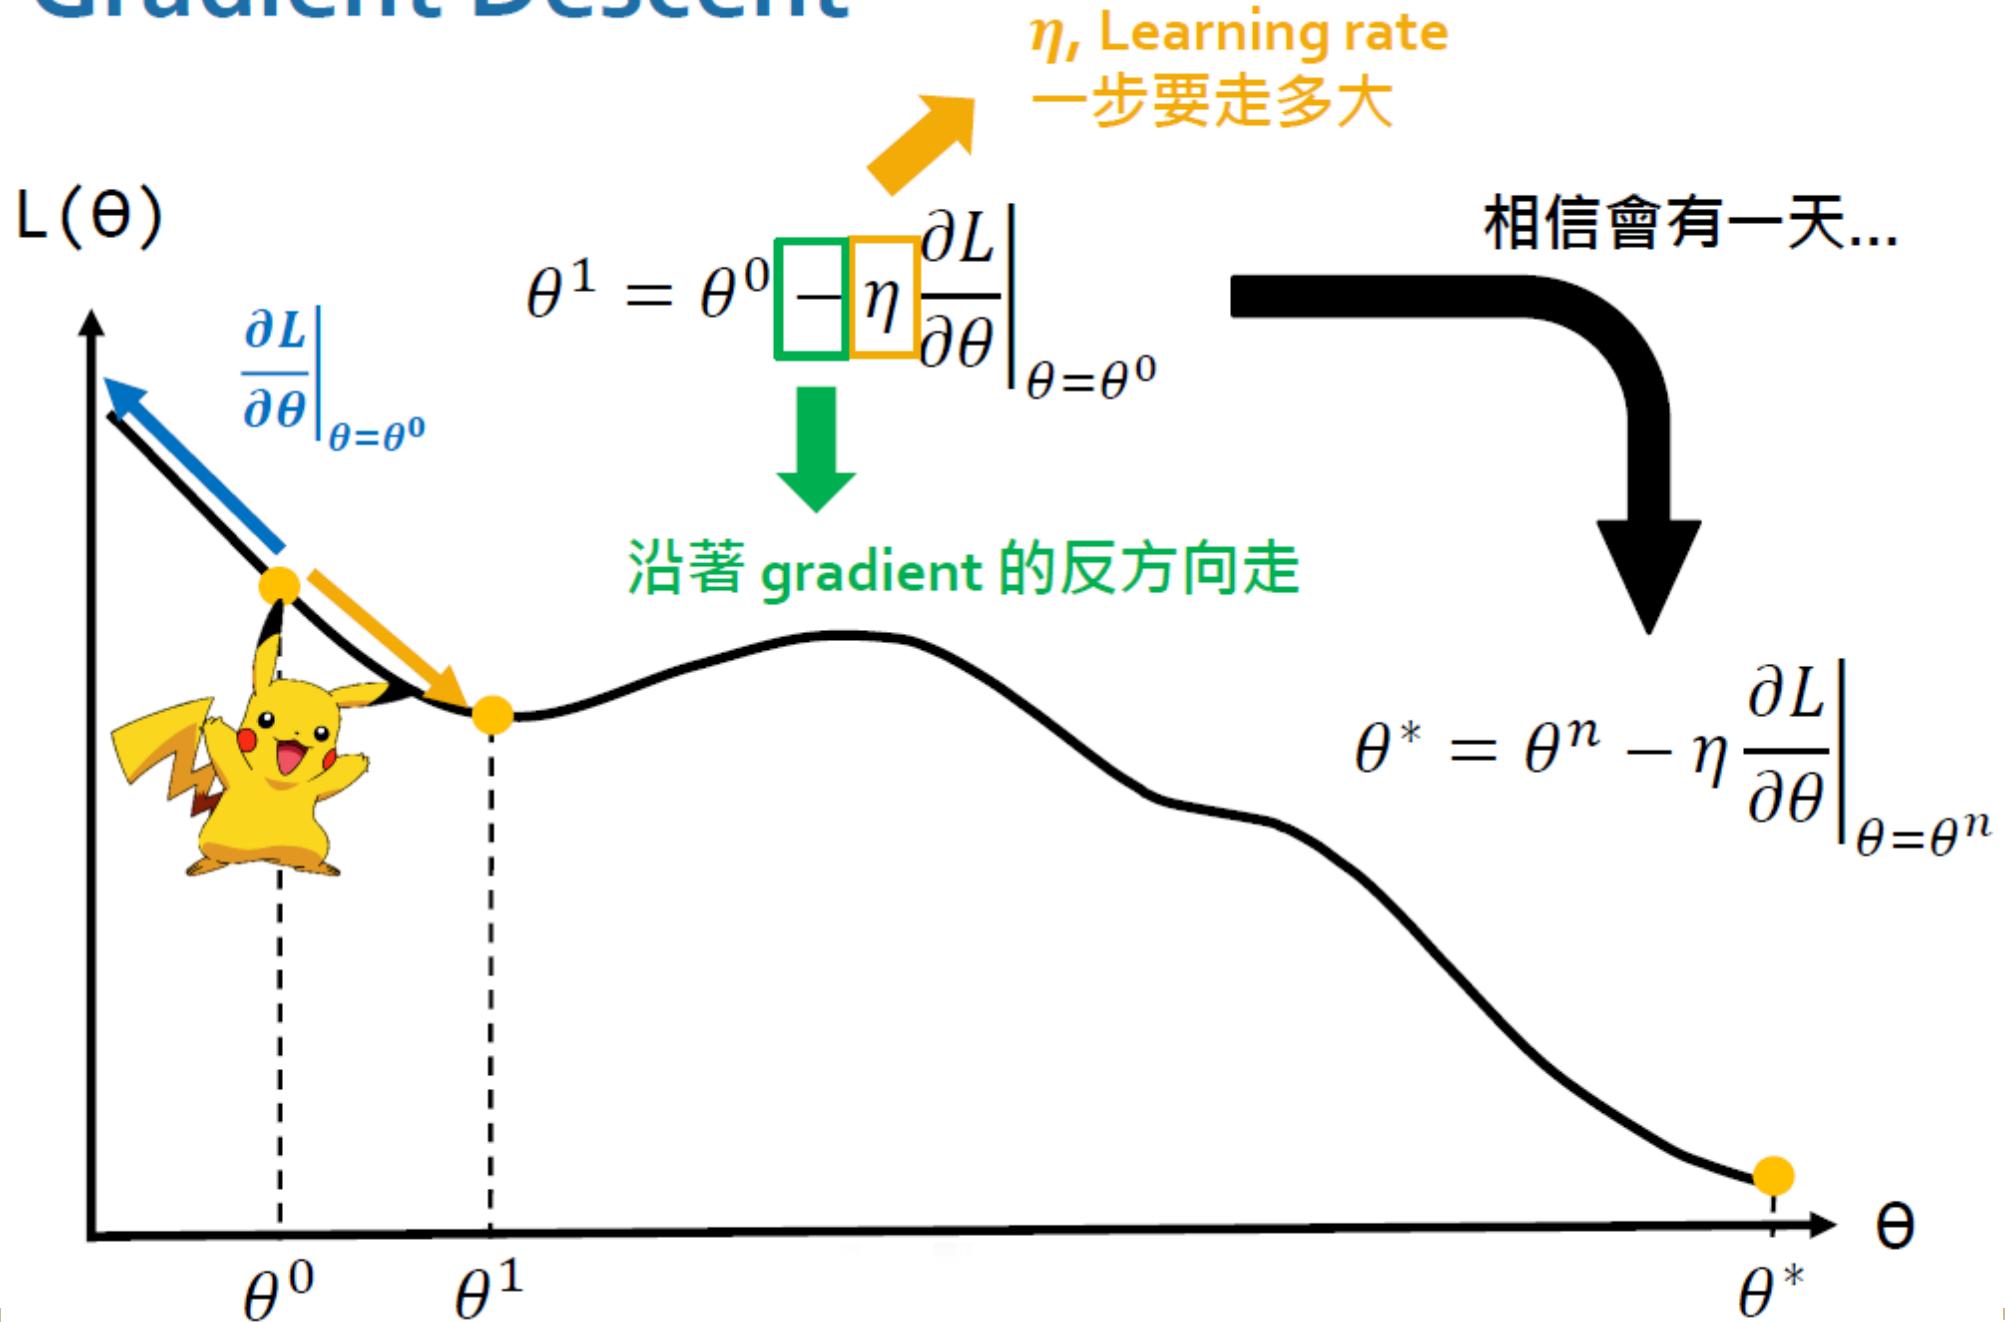


# Outline

- Gradient descent variants
  - Batch gradient descent/Stochastic gradient descent/ **Mini-batch gradient descent**
- Challenges
- Gradient descent optimization algorithms
  - Momentum/Nesterov accelerated gradient
  - Adagrad/Adadelta/RMSprop/
  - **Adam**/AdaMax//Nadam
  - Visualization of algorithms
  - Which optimizer to choose?

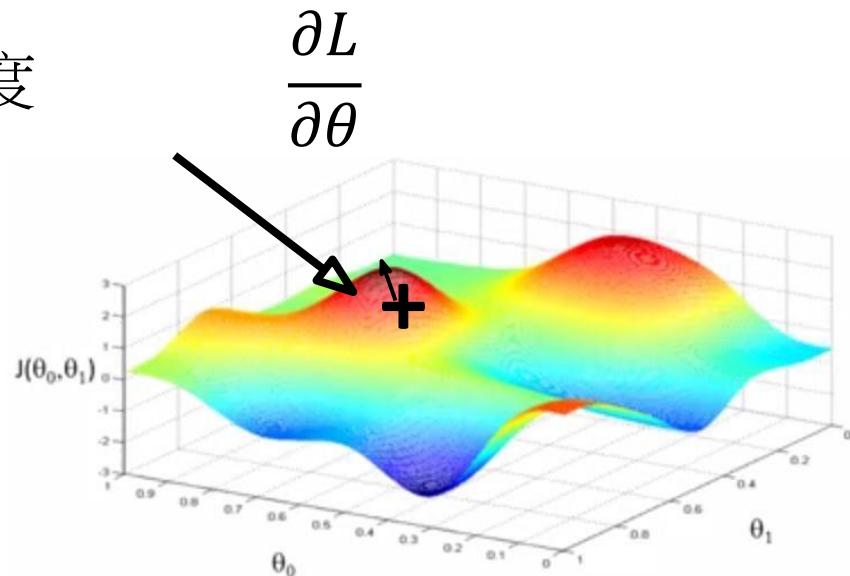
# INTRODUCTION

# Gradient Descent

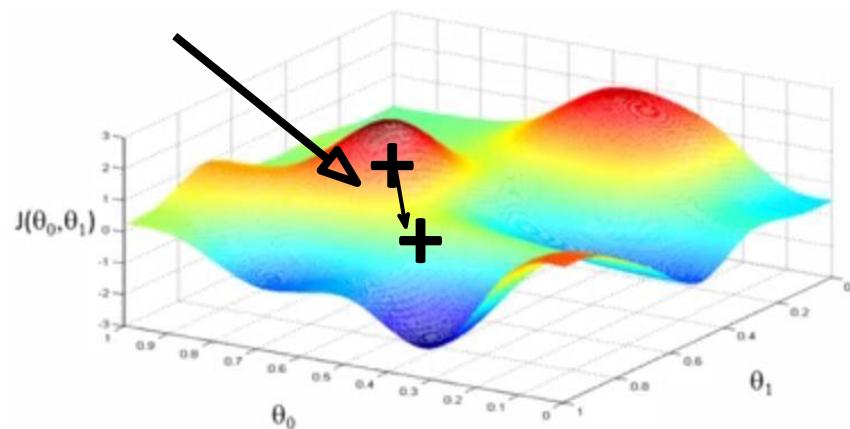


# Gradient Descent (梯度下降法)

計算梯度



沿著 gradient 的反方向走



Initialize  $\theta$  randomly

For N Epochs

- For each training example  $(x, y)$ :
  - Compute Loss Gradient:  $\frac{\partial L}{\partial \theta}$
  - Update  $\theta$  with update rule:

$$\theta = \theta - \eta \frac{\partial L}{\partial \theta}$$

沿著 gradient 的反方向走

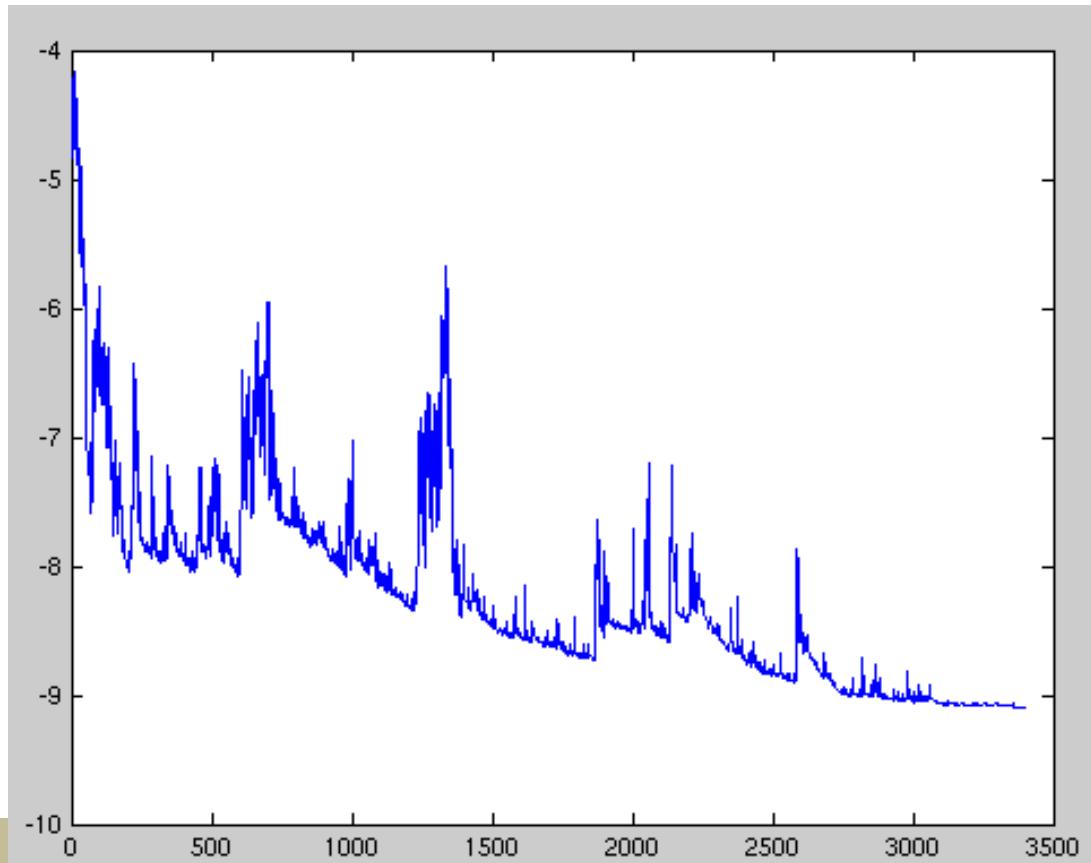
Learning rate

# Three Variants

- Batch gradient descent
  - Vanilla, original
  - Compute loss over whole training dataset at a time
- Stochastic gradient descent
  - Loss measured per example & label
- Mini-batch gradient descent
  - “Best of both worlds”
  - Loss computed over subset of training dataset

# Stochastic Gradient Descent

- SGD performs frequent updates
  - high variance
  - Causes objective function to fluctuate heavily

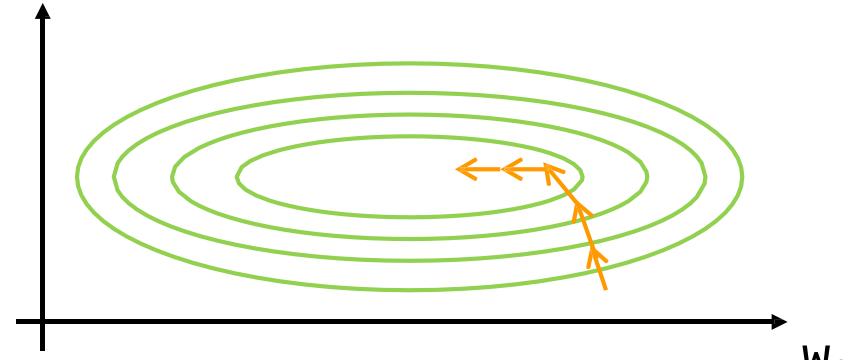


# Stochastic Gradient Descent

- Other advantages of SGD vs. batch?
  - enables it to jump to new and potentially better local minima
- Side-effect of SGD's approach to finding minima?
  - Complicates convergence: SGD will keep overshooting
- Solution for this?
  - Slowly decrease learning rate
- SGD shows the **same convergence behaviour** as batch gradient
  - Almost certainly converging to a **local minimum for non-convex optimisation**
  - Almost certainly converging to a **global minimum for convex optimisation**

# Challenges of Vanilla mini-batch gradient descent

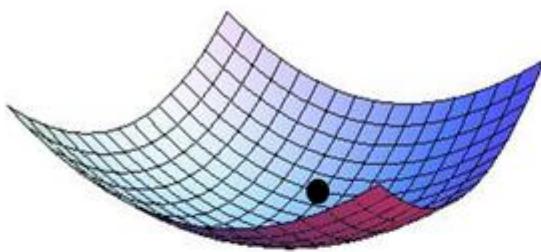
- No guarantee good convergence
- Choosing a proper learning rate can be difficult
  - Too small: very slow convergence
  - Too high:
    - Damages convergence
    - Loss rate can fluctuate or even diverge
- Learning rate decay should be defined in advance (time or threshold)
  - unable to adapt to a dataset's characteristics
- Same learning rate for all parameters updates
  - Not all features have the same variance
  - Bad for sparse data and features with different frequencies (e.g. need larger update for rarely occurring features)



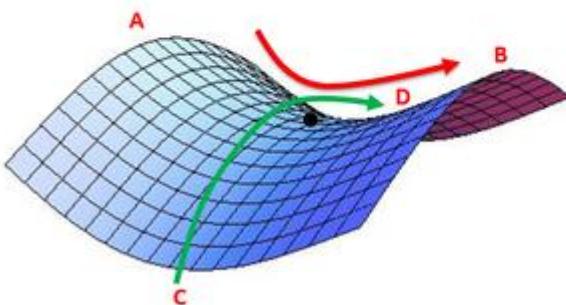
- **Saddle points are a challenge**
  - Neural nets tend to have many local minima
  - These tend to meet in saddle points
  - Equates to a plateau of the same error
  - Hard to escape using gradient-based method, as  $\text{gradient} \rightarrow 0$  in all dimensions

$$\theta = \theta - \eta \frac{\partial L}{\partial \theta}$$

Solutions: adaptive learning rate strategy  
+ momentum

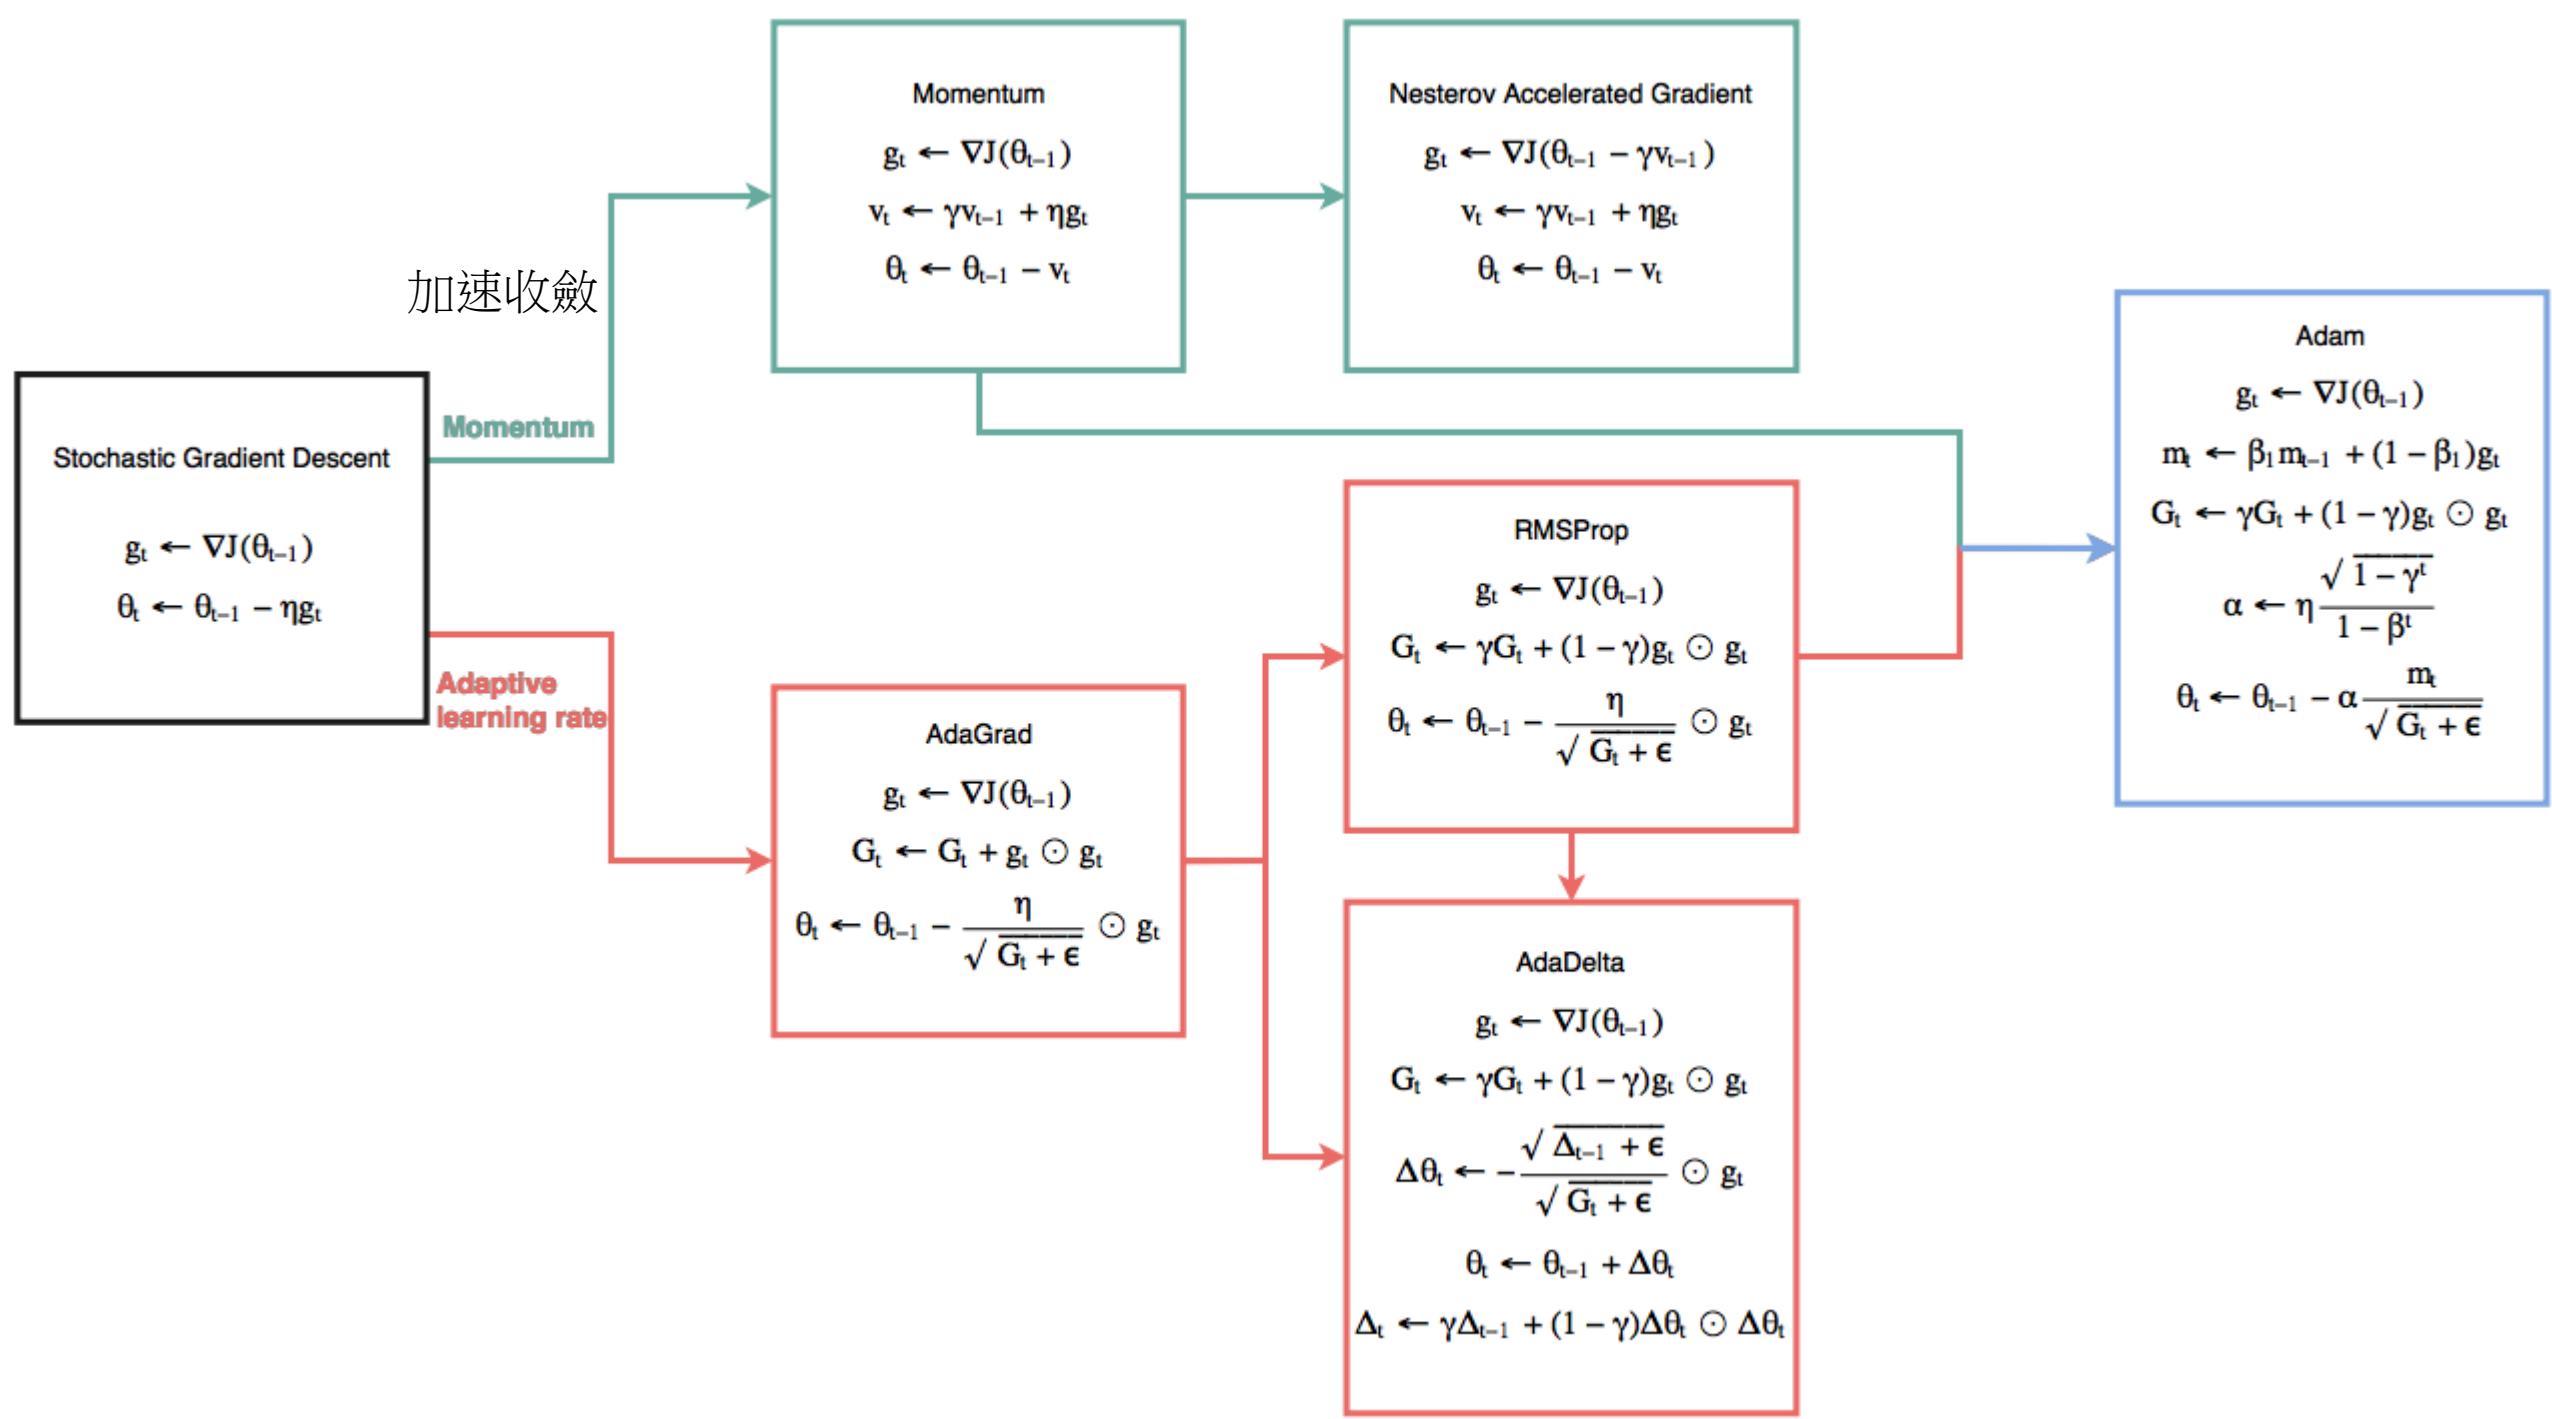


Minima- Black dot placed on the PES shows a minimum energy point. Note how a PES resembles a well around the minimum point.



Saddle Point- Black dot placed on the PES shows a minima along path A-B and a maxima along path C-D. It represents a transition state along path C-D which, in this case, is the reaction coordinate.

# FROM SGD TO ADAM AND ITS VARIANTS



# Gradient descent optimization

- Momentum based improvement
  - Momentum
  - Nesterov Momentum
- Per-parameter adaptive learning rate
  - Adagrad
  - Adadelta
  - RMSprop
- Combination of above two
  - ADAM

SGD

$$\mathbf{g}_t \leftarrow \nabla J_i(\boldsymbol{\theta}_{t-1})$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \eta \mathbf{g}_t$$

Minibatch SGD

$$\mathbf{g}_t \leftarrow \frac{1}{n} \sum_i^n \nabla J_i(\boldsymbol{\theta}_{t-1})$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \eta \mathbf{g}_t$$



Momentum

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$\mathbf{v}_t \leftarrow \underline{\gamma \mathbf{v}_{t-1}} + \eta \mathbf{g}_t$$

$$\boldsymbol{\theta}_t \leftarrow \underline{\boldsymbol{\theta}_{t-1}} - \mathbf{v}_t$$



Nesterov Momentum

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1} - \underline{\gamma \mathbf{v}_{t-1}})$$

$$\mathbf{v}_t \leftarrow \gamma \mathbf{v}_{t-1} + \eta \mathbf{g}_t$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \mathbf{v}_t$$

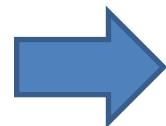
梯度是過去梯度的加權累積

Keep accumulated  
momentum to escape  
local minimumCalculate gradient based  
on predicted next valueWeight 先用預測值假裝更新，  
再算梯度，做真的更新，會比較  
準

SGD

$$\mathbf{g}_t \leftarrow \nabla J_i(\boldsymbol{\theta}_{t-1})$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \eta \mathbf{g}_t$$

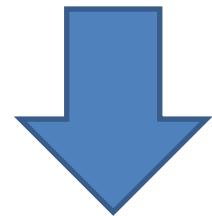


AdaGrad

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$G_t \leftarrow G_t + \mathbf{g}_t \odot \mathbf{g}_t$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$$



Use  $\sum_{i=1}^t \mathbf{g}_{i,j}^2$  to judge how often  
Weight<sub>i</sub> have been updated  
(Arithmetic average)

Larger G<sub>t</sub> → freq. update  
→ smaller learning rate

Small G<sub>t</sub> → seldom update  
→ larger learning rate

Note. G<sub>t</sub> is monotonically increased  
→ learning rate => 0

RMSProp

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$G_t \leftarrow \frac{\gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t}{\eta}$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$$

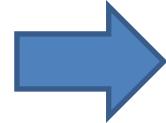
Degraded G<sub>t</sub> effect (Exponential average)  
Avoid learning rate becomes zero

## RMSProp

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$G_t \leftarrow \underline{\gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t}$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$$



## AdaDelta

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$G_t \leftarrow \gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t$$

$$\Delta \boldsymbol{\theta}_t \leftarrow -\frac{\sqrt{\Delta_{t-1} + \epsilon}}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} + \Delta \boldsymbol{\theta}_t$$

$$\underline{\Delta_t \leftarrow \gamma \Delta_{t-1} + (1 - \gamma) \Delta \boldsymbol{\theta}_t \odot \Delta \boldsymbol{\theta}_t}$$

估計 learning rate

Automated learning rate  
instead of fixed one

利用之前的步長們估計下一步的  
步長

可能問題;  $G_t$  是固定時間窗口內的累積，隨著時間窗口的變化，遇到的數據可能發生巨變，使得  $G_t$  可能會時大時小，不是單調變化。這就可能在訓練後期引起學習率的震盪，導致模型無法收斂

## Momentum

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$\mathbf{v}_t \leftarrow \underline{\gamma \mathbf{v}_{t-1}} + \eta \mathbf{g}_t$$

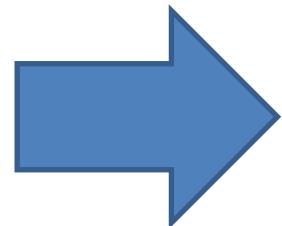
$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \mathbf{v}_t$$

## RMSProp

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1})$$

$$G_t \leftarrow \underline{\gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t}$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot \mathbf{g}_t$$



## ADAM

$$\mathbf{g}_t \leftarrow \nabla J(\boldsymbol{\theta}_{t-1}) \quad \text{Adam做一階/二階動量估計}$$

$$\mathbf{m}_t \leftarrow \underline{\beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t} \quad \text{Momentum}$$

$$\underline{G_t \leftarrow \gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t} \quad \text{RMSProp}$$

$$\alpha \leftarrow \eta \frac{\sqrt{1 - \gamma^t}}{1 - \beta^t}$$

$$\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} - \alpha \frac{\mathbf{m}_t}{\sqrt{G_t + \epsilon}}$$

Bias correction term

# AdamW (Adam with decoupled weight decay)

**Algorithm 2** Adam with L<sub>2</sub> regularization and Adam with decoupled weight decay (AdamW)

- ```

1: given  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ ,  $\epsilon = 10^{-8}$ ,  $\lambda \in \mathbb{R}$ 
2: initialize time step  $t \leftarrow 0$ , parameter vector  $\theta_{t=0} \in \mathbb{R}^n$ , first moment vector  $m_{t=0} \leftarrow \theta$ , second moment
   vector  $v_{t=0} \leftarrow \theta$ , schedule multiplier  $\eta_{t=0} \in \mathbb{R}$ 
3: repeat
4:    $t \leftarrow t + 1$ 
5:    $\nabla f_t(\theta_{t-1}) \leftarrow \text{SelectBatch}(\theta_{t-1})$                                  $\triangleright$  select batch and return the corresponding gradient
6:    $\mathbf{g}_t \leftarrow \nabla f_t(\theta_{t-1}) + \lambda \theta_{t-1}$ 
7:    $\mathbf{m}_t \leftarrow \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$                                  $\triangleright$  here and below all operations are element-wise
8:    $\mathbf{v}_t \leftarrow \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2$ 
9:    $\hat{\mathbf{m}}_t \leftarrow \mathbf{m}_t / (1 - \beta_1^t)$  $\triangleright \beta_1$  is taken to the power of  $t$ 
10:   $\hat{\mathbf{v}}_t \leftarrow \mathbf{v}_t / (1 - \beta_2^t)$  $\triangleright \beta_2$  is taken to the power of  $t$ 
11:   $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$                                  $\triangleright$  can be fixed, decay, or also be used for warm restarts
12:   $\theta_t \leftarrow \theta_{t-1} - \eta_t \left( \alpha \hat{\mathbf{m}}_t / (\sqrt{\hat{\mathbf{v}}_t} + \epsilon) + \lambda \theta_{t-1} \right)$ 
13: until stopping criterion is met
14: return optimized parameters  $\theta_t$ 

```

# 為什麼要「Decoupled Weight Decay」

## AdamW 介紹

- 避免overfitting，兩種常用的正規化策略
  - L2 regularization in loss function
    - `final_loss = loss + wd * all_weights.pow(2).sum() / 2`
    - 沒有weight decay,  $w = w - lr * \nabla L$
    - 對SGD,  $w = w - lr * \nabla L_{final} = w - lr * (\nabla L + wd * w) = w(1 - wd) - lr * \nabla L$
  - Weight decay
    - $w = w - lr * w.grad - lr * wd * w = (1-wd)*w - lr * w.grad$

$$\mathbf{g}_t \leftarrow \nabla J(\theta_{t-1})$$

$$\mathbf{m}_t \leftarrow \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$$

$$G_t \leftarrow \gamma G_t + (1 - \gamma) \mathbf{g}_t \odot \mathbf{g}_t$$

$$\alpha \leftarrow \eta \frac{\sqrt{1 - \gamma^t}}{1 - \beta^t}$$

$$\theta_t \leftarrow \theta_{t-1} - \alpha \frac{\mathbf{m}_t}{\sqrt{G_t + \epsilon}}$$

$$\nabla L_{final} = (\nabla L + \lambda * \theta)$$

用 L2 regularization  $g_t = g_{t-1} + \lambda \theta$

用 weight decay  $\theta = \theta - \alpha \nabla L_{final} = (1 - \lambda) \theta - \alpha \nabla L$

對 **SGD** 這等價於 **weight decay** (每步把  $\theta$  乘上  $1 - \lambda$ )。

但對 **Adam**，因為會除以  $\text{sqrt}(G_t)$  也被自適應縮放，導致每一維的衰減「不再等比例」，正則化力道被梯度統計扭曲，使得「L2 ≠ weight decay」

**AdamW** 的關鍵：把衰減從梯度路徑中拿出來獨立處理  
這樣衰減量與  $v^t$  無關，回到「真正的 weight decay」  
(各維等比例收縮)，正則效果更可控、與學習率/排程的互動也更直觀。

$$\theta_{t+1} = \theta_t - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon} \underbrace{- \alpha \cdot \text{wd} \cdot \theta_t}_{\text{decoupled weight decay}}$$

# Lion (Evolved Sign Momentum)

- Algorithm searched optimizer
- Compared to Adam
  - Lower **memory** usage (only keep track of the momentum)
- Different from adaptive optimizers,
  - its update has the same magnitude for each parameter calculated through the **sign** operation
- Performance
  - Up to 2% increase for ViT
  - Reduce training time by up to 2.3X for diffusion models

$$\text{Lion} := \begin{cases} \mathbf{u}_t = \underline{\text{sign}}\left(\beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t\right) \\ \boldsymbol{\theta}_t = \boldsymbol{\theta}_{t-1} - \eta_t (\mathbf{u}_t + \lambda_t \boldsymbol{\theta}_{t-1}) \\ \mathbf{m}_t = \beta_2 \mathbf{m}_{t-1} + (1 - \beta_2) \mathbf{g}_t \end{cases}$$

Momentum 的更新放在weight  
更新之後

$\beta_1=0.9, \beta=0.99$

Lion的更新量 $\mathbf{u}$ 每個分量的絕對值都是1，這通常比AdamW要大，所以學習率要縮小10倍以上

$$\text{AdamW} := \begin{cases} \mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t \\ \mathbf{v}_t = \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2 \\ \hat{\mathbf{m}}_t = \mathbf{m}_t / (1 - \beta_1^t) \\ \hat{\mathbf{v}}_t = \mathbf{v}_t / (1 - \beta_2^t) \\ \mathbf{u}_t = \hat{\mathbf{m}}_t / (\sqrt{\hat{\mathbf{v}}_t} + \epsilon) \\ \boldsymbol{\theta}_t = \boldsymbol{\theta}_{t-1} - \eta_t (\mathbf{u}_t + \lambda_t \boldsymbol{\theta}_{t-1}) \end{cases}$$

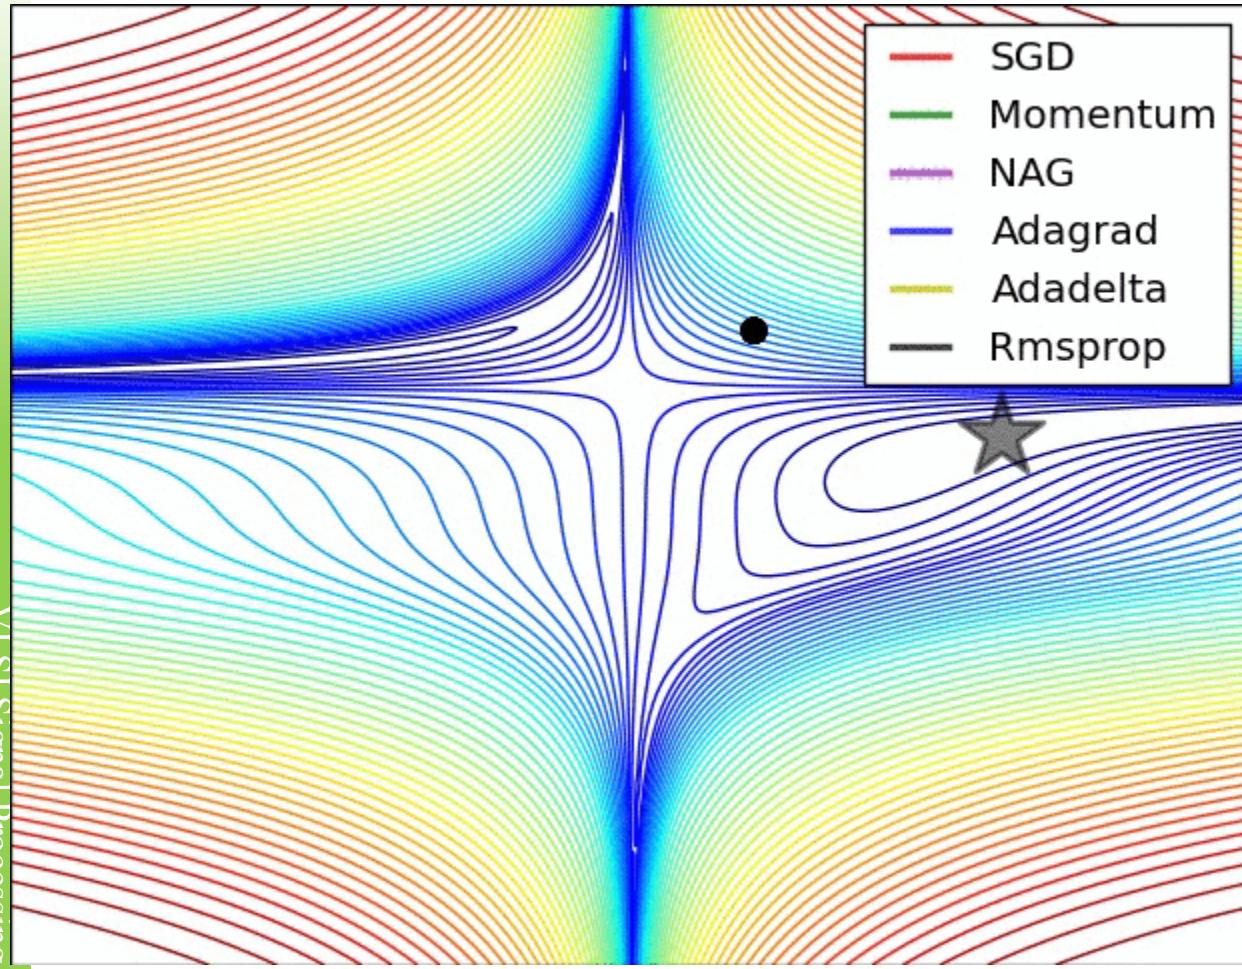
Table 7: The performance of various optimizers to train ViT-S/16 and ViT-B/16 on ImageNet (with RandAug and Mixup). Lion is still the best performing one, and there is no clear winner amongst the baselines.

| Model    | Task     | AdamW | RAdam | NAdam | AdaBelief | AMSGrad | Ablation <sub>0.9</sub> | Ablation <sub>0.99</sub> | Lion         |
|----------|----------|-------|-------|-------|-----------|---------|-------------------------|--------------------------|--------------|
| ViT-S/16 | ImageNet | 78.89 | 78.59 | 78.91 | 78.71     | 79.01   | 78.23                   | 78.19                    | <b>79.46</b> |
|          | ReaL     | 84.61 | 84.47 | 84.62 | 84.56     | 85.01   | 84.28                   | 84.17                    | <b>85.25</b> |
|          | V2       | 66.73 | 66.39 | 66.02 | 66.35     | 66.82   | 66.13                   | 65.96                    | <b>67.68</b> |
| ViT-B/16 | ImageNet | 80.12 | 80.26 | 80.32 | 80.29     | 79.85   | 79.54                   | 79.90                    | <b>80.77</b> |
|          | ReaL     | 85.46 | 85.45 | 85.44 | 85.48     | 85.16   | 85.10                   | 85.36                    | <b>86.15</b> |
|          | V2       | 68.14 | 67.76 | 68.46 | 68.19     | 68.48   | 68.07                   | 68.20                    | <b>69.19</b> |

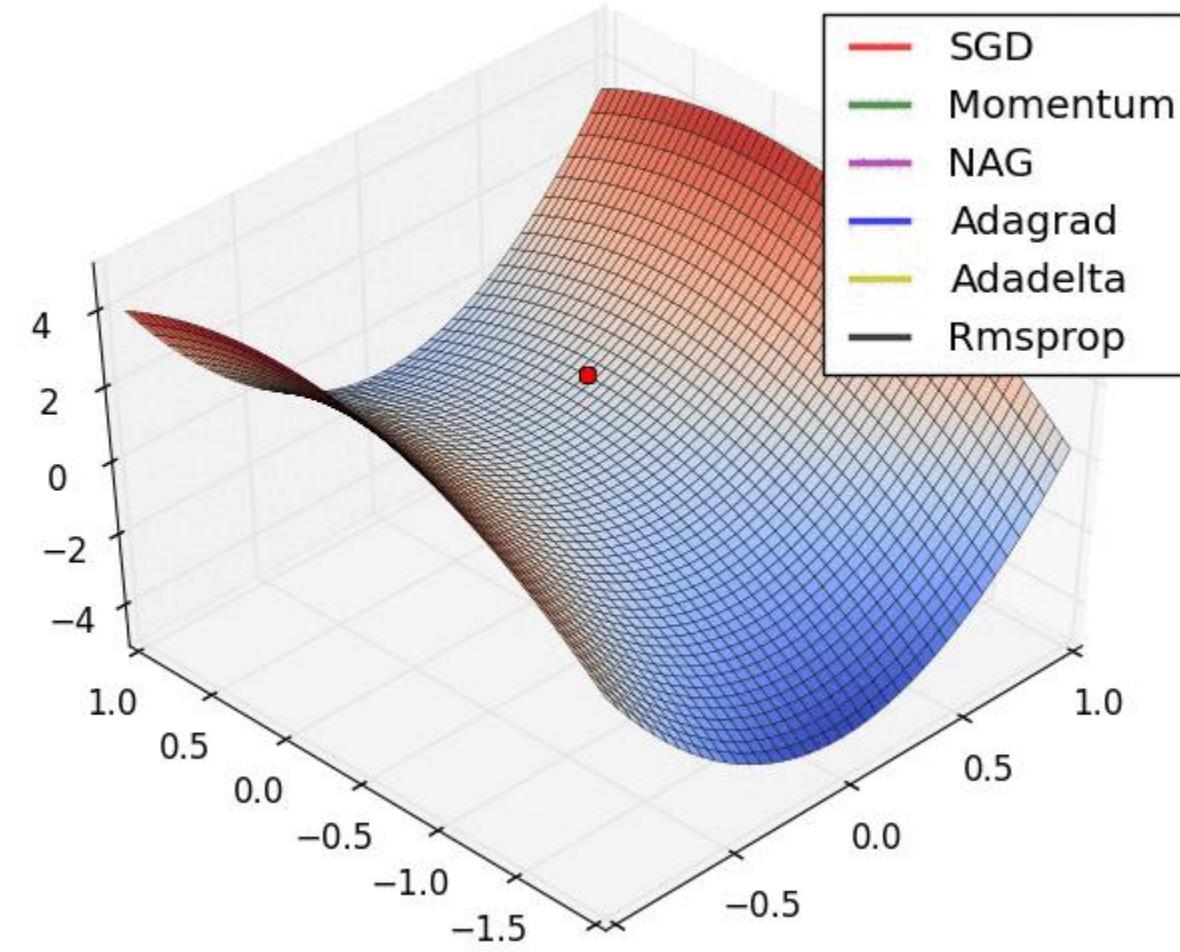
Table 5: One-shot evaluation averaged over three NLG and 21 NLU tasks. The results of GPT-3 (Brown et al., 2020) and PaLM (Chowdhery et al., 2022) are included for reference. The LLMs trained by Lion have better in-context learning ability. See Table 11 (in the Appendix) for detailed results on all tasks.

| Task    | 1.1B      |             | 2.1B      |             | 7.5B      |             | 6.7B  |      | 8B |  |
|---------|-----------|-------------|-----------|-------------|-----------|-------------|-------|------|----|--|
|         | Adafactor | Lion        | Adafactor | Lion        | Adafactor | Lion        | GPT-3 | PaLM |    |  |
| #Tokens | 300B      |             |           |             |           |             | 300B  | 780B |    |  |
| Avg NLG | 11.1      | <b>12.1</b> | 15.6      | <b>16.5</b> | 24.1      | <b>24.7</b> | 23.1  | 23.9 |    |  |
| Avg NLU | 53.2      | <b>53.9</b> | 56.8      | <b>57.4</b> | 61.3      | <b>61.7</b> | 58.5  | 59.4 |    |  |

# Visualization



SGD optimization on loss surface contours



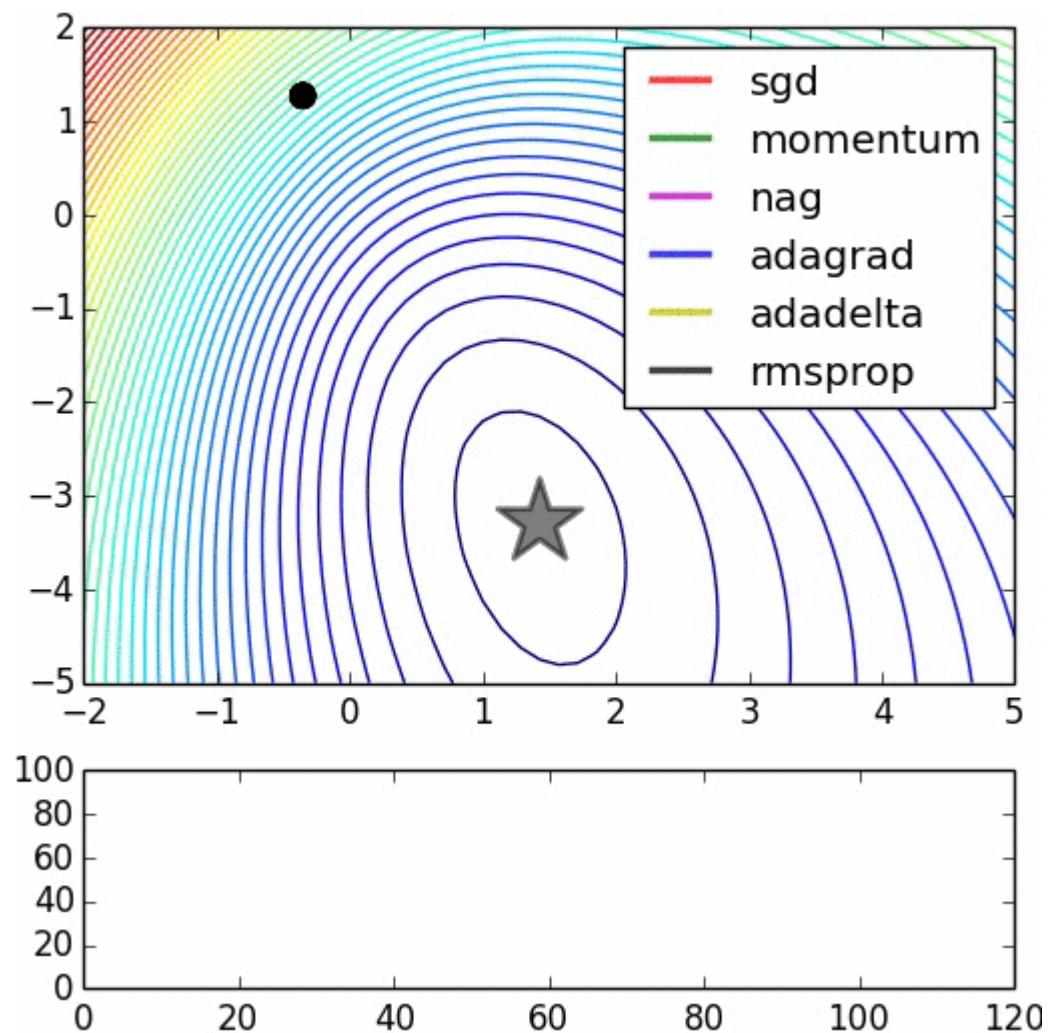
SGD optimization on saddle point

- Behaviour on the contours of a loss surface
  - Adagrad, Adadelta, and RMSprop almost immediately head off in the right direction and converge similarly fast,
  - Momentum and NAG are led off-track, evoking the image of a ball rolling down the hill.
  - NAG, however, is quickly able to correct its course due to its increased responsiveness by looking ahead and heads to the minimum
  - Notice the "overshooting" behavior of momentum-based methods, which make the optimization look like a ball rolling down the hill

# Behaviour on the saddle points

- SGD, Momentum, and NAG find it difficult to break symmetry, although the two latter eventually manage to escape the saddle point
- Adagrad, RMSprop, and Adadelta quickly head down the negative slope.

# Learning rate v.s. Gradient Optimization



Appendix

**IMPROVE ADAM  
WE WILL ONLY GO THROUGH PART OF THIS**

# Improvement

- +weight decay or L2 regularization
  - ADAMW: use weight decay
- Initial warmup for training stability
  - RADAM: rectified ADAM, 避免起始model 不穩定，根據方差分散度，動態地打開或者關閉自適應學習率，並且提供了一種不需要可調參數學習率預熱的方法。
    - 對起始warmup 有用，後期類似ADAM
  - AdaMod: ADAM + long term memory of learning rate (long term average of the adaptive learning rates)
- 減少超參數調整
  - Lookahead optimizer: fast and slow weight
  - Ranger: RADAM + Lookahead
- Better convergence
  - Gradient centralization: zero mean gradient
  - diffGrad: 控制收斂時梯度，避免overshooting

# AdamW: Adam + weight decay regularization

- Weight decay or L2 regularization?

```
final_loss = loss + wd * all_weights.pow(2).sum() / 2  
  
w = w - lr * w.grad - lr * wd * w
```

- These two are equivalent for vanilla SGD,
- But become different with momentum

```
//L2 regularization update  
moving_avg = alpha * moving_avg + (1-alpha) * (w.grad + wd*w)  
//weight decay's update  
moving_avg = alpha * moving_avg + (1-alpha) * w.grad  
w = w - lr * moving_avg - lr * wd * w
```

- Suggest to use weight decay version (ADAMW)

# Adafactor (Adaptive Learning Rates with Factorization)

針對NLP (LLM) 應用，省記  
憶體用量

$$\begin{cases} g_t = \nabla_{\theta} L(\theta_{t-1}) \\ m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\ \hat{m}_t = m_t / (1 - \beta_1^t) \\ \hat{v}_t = v_t / (1 - \beta_2^t) \\ \theta_t = \theta_{t-1} - \alpha_t \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) \end{cases}$$

再加入正確的weight decay 和layer adaptive

對NLP，adaptive learning 比較重要，  
拋棄動量=>類RMSprop，變數直接減  
少了一半

$$\begin{cases} g_t = \nabla_{\theta} L(\theta_{t-1}) \\ v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\ \hat{v}_t = v_t / (1 - \beta_2^t) \\ \theta_t = \theta_{t-1} - \alpha_t g_t / \sqrt{\hat{v}_t + \epsilon} \end{cases}$$

$$\begin{cases} g_{i,j;t} = \nabla_{\theta} L(\theta_{i,j;t-1}) \\ \hat{\beta}_{2,t} = 1 - t^{-c} \\ v_{i;t}^{(r)} = \hat{\beta}_{2,t} v_{t-1;i}^{(r)} + (1 - \hat{\beta}_{2,t}) \sum_j (g_{i,j;t}^2 + \epsilon_1) \\ v_{j;t}^{(c)} = \hat{\beta}_{2,t} v_{t-1;j}^{(c)} + (1 - \hat{\beta}_{2,t}) \sum_i (g_{i,j;t}^2 + \epsilon_1) \\ \hat{v}_{i,j;t} = v_{i;t}^{(r)} v_{j;t}^{(c)} / \sum_j v_{j;t}^{(c)} \end{cases}$$

$$\begin{cases} u_t = g_t / \sqrt{\hat{v}_t} \\ \hat{u}_t = u_t / \max(1, RMS(u_t)/d) \times \max(\epsilon_2, RMS(\theta_{t-1})) \\ \theta_t = \theta_{t-1} - \alpha_t \hat{u}_t \end{cases}$$

把變數v的參數量再壓縮。用矩陣的低  
秩分解

$$\begin{cases} g_{i,j;t} = \nabla_{\theta} L(\theta_{i,j;t-1}) \\ v_{i;t}^{(r)} = \beta_2 v_{t-1;i}^{(r)} + (1 - \beta_2) \sum_j (g_{i,j;t}^2 + \epsilon) \\ v_{j;t}^{(c)} = \beta_2 v_{t-1;j}^{(c)} + (1 - \beta_2) \sum_i (g_{i,j;t}^2 + \epsilon) \\ v_{i,j;t} = v_{i;t}^{(r)} v_{j;t}^{(c)} / \sum_j v_{j;t}^{(c)} \\ \hat{v}_t = v_t / (1 - \beta_2^t) \\ \theta_t = \theta_{t-1} - \alpha_t g_t / \sqrt{\hat{v}_t} \end{cases}$$

# Lion: EvoLved Sign Momentum

---

**Algorithm 1** AdamW Optimizer
 

---

```

given  $\beta_1, \beta_2, \epsilon, \lambda, \eta, f$ 
initialize  $\theta_0, m_0 \leftarrow 0, v_0 \leftarrow 0, t \leftarrow 0$ 
while  $\theta_t$  not converged do
   $t \leftarrow t + 1$ 
   $g_t \leftarrow \nabla_{\theta} f(\theta_{t-1})$ 
  update EMA of  $g_t$  and  $g_t^2$ 
   $m_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$ 
   $v_t \leftarrow \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$  ← Lion 省了二階
  bias correction
   $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ 
   $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  ← Lion 省了二階
  update model parameters
   $\theta_t \leftarrow \theta_{t-1} - \eta_t (\hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon) + \lambda \theta_{t-1})$ 
end while
return  $\theta_t$ 
  
```

---



---

**Algorithm 2** Lion Optimizer (ours)
 

---

```

given  $\beta_1, \beta_2, \lambda, \eta, f$ 
initialize  $\theta_0, m_0 \leftarrow 0$ 
while  $\theta_t$  not converged do
   $g_t \leftarrow \nabla_{\theta} f(\theta_{t-1})$ 
  update model parameters
   $c_t \leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t$ 
   $\theta_t \leftarrow \theta_{t-1} - \eta_t (\text{sign}(c_t) + \lambda \theta_{t-1})$ 
  update EMA of  $g_t$ 
   $m_t \leftarrow \beta_2 m_{t-1} + (1 - \beta_2) g_t$ 
end while
return  $\theta_t$ 
  
```

---

與 AdamW 和各種自適應優化器需要同時保存一階和二階矩相比，Lion 只需要動量，將額外的記憶體佔用減半。這在訓練大型模型和/或大批量時很有用。小 batch\_size（小於 64）的時候效果不如 AdamW

# Sophia: A Scalable Stochastic Second-order Optimizer for Large Model Pre-training

## Algorithm 1 Hutchinson( $\theta$ )

- 1: **Input:** parameter  $\theta$ .
- 2: Compute mini-batch loss  $L(\theta)$ .
- 3: Draw  $u$  from  $\mathcal{N}(0, \mathbf{I}_d)$ .
- 4: **return**  $u \odot \nabla(\langle \nabla L(\theta), u \rangle)$ .

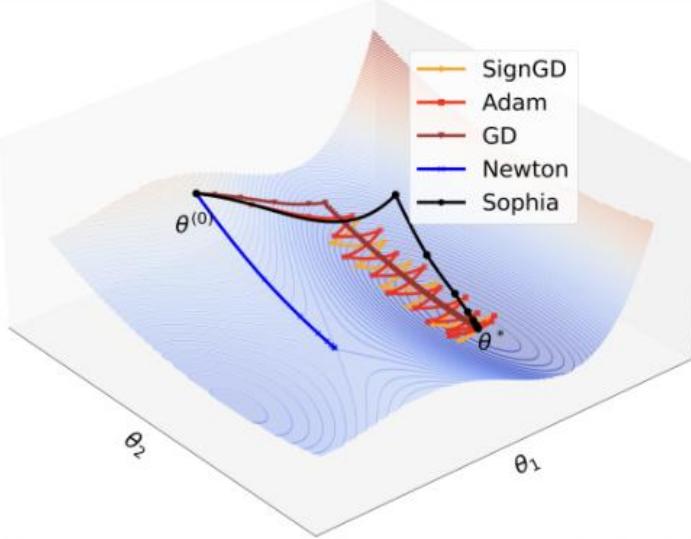


Figure 2: The motivating toy example.  $\theta_{[1]}$  is the sharp dimension and  $\theta_{[2]}$  is the flat dimension. GD's learning rate is limited by the sharpness in  $\theta_1$ , and makes slow progress along  $\theta_{[2]}$ . Adam and SignGD bounce along  $\theta_{[1]}$  while making slow progress along  $\theta_{[2]}$ . Vanilla Newton's method converges to a saddle point. Sophia makes fast progress in both dimensions and converges to the minimum with a few steps.

## Algorithm 2 Gauss-Newton-Bartlett( $\theta$ )

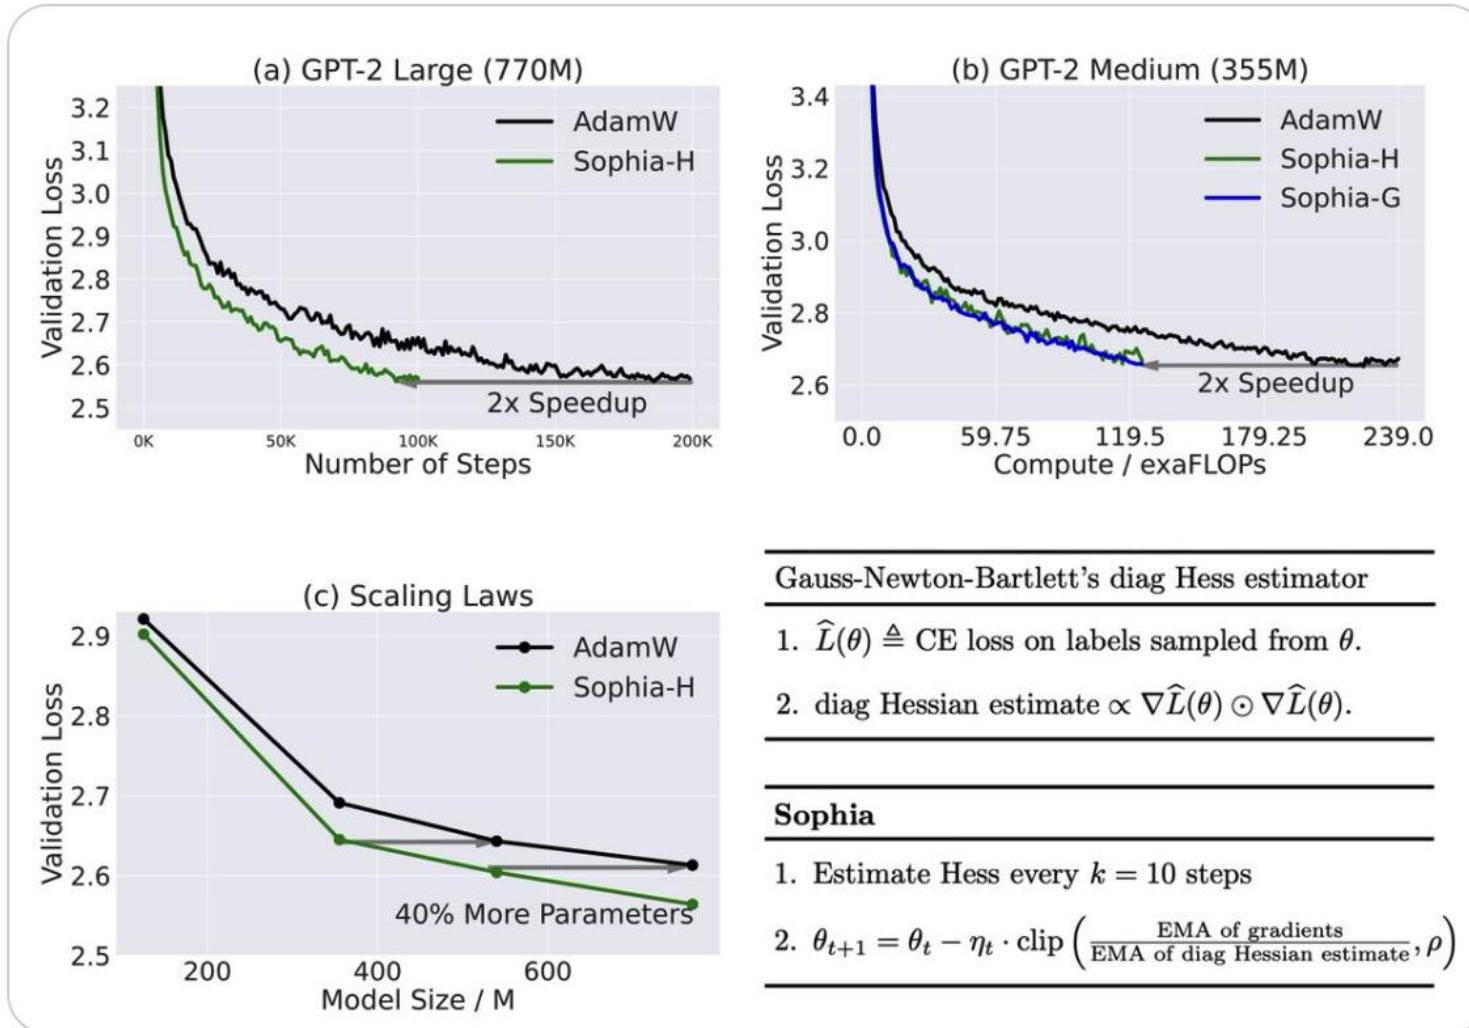
- 1: **Input:** parameter  $\theta$ .
- 2: Draw a mini-batch of input  $\{x_b\}_{b=1}^B$ .
- 3: Compute logits on the mini-batch:  $\{f(\theta, x_b)\}_{b=1}^B$ .
- 4: Sample  $\hat{y}_b \sim \text{softmax}(f(\theta, x_b)), \forall b \in [B]$ .
- 5: Calculate  $\hat{g} = \nabla(1/B \sum \ell(f(\theta, x_b), \hat{y}_b))$ .
- 6: **return**  $B \cdot \hat{g} \odot \hat{g}$ .

## Algorithm 3 Sophia

- 1: **Input:**  $\theta_1$ , learning rate  $\{\eta_t\}_{t=1}^T$ , hyperparameters  $\lambda, \gamma, \beta_1, \beta_2, \epsilon$ , and estimator choice  $\text{Estimator} \in \{\text{Hutchinson}, \text{Gauss-Newton-Bartlett}\}$
- 2: Set  $m_0 = 0, v_0 = 0, h_{1-k} = 0$
- 3: **for**  $t = 1$  to  $T$  **do**
- 4:   Compute minibatch loss  $L_t(\theta_t)$ .
- 5:   Compute  $g_t = \nabla L_t(\theta_t)$ .
- 6:    $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
- 7:   **if**  $t \bmod k = 1$  **then**
- 8:     Compute  $\hat{h}_t = \text{Estimator}(\theta_t)$ .
- 9:      $h_t = \beta_2 h_{t-k} + (1 - \beta_2) \hat{h}_t$
- 10:   **else**
- 11:      $h_t = h_{t-1}$
- 12:   **end if**
- 13:    $\theta_t = \theta_t - \eta_t \lambda \theta_t$  (weight decay)
- 14:    $\theta_{t+1} = \theta_t - \eta_t \cdot \text{clip}(m_t / \max\{\gamma \cdot h_t, \epsilon\}, 1)$

Sophia 優化器使用隨機估計值作為 Hessian 矩陣對角線的 pre-conditioner，並採用剪切 (clipping) 機制來控制最壞情況下的參數大小更新。在像 GPT-2 這樣的預訓練語言模型上，Sophia 與 Adam 相比，在減少了 50% step 數量的情況下實現了相同的驗證預訓練損失。

# Sophia: A Scalable Stochastic Second-order Optimizer for Language Model Pre-training




---

### Gauss-Newton-Bartlett's diag Hess estimator

---

1.  $\hat{L}(\theta) \triangleq$  CE loss on labels sampled from  $\theta$ .
  2. diag Hessian estimate  $\propto \nabla \hat{L}(\theta) \odot \nabla \hat{L}(\theta)$ .
- 

---

### Sophia

---

1. Estimate Hess every  $k = 10$  steps
  2. 
$$\theta_{t+1} = \theta_t - \eta_t \cdot \text{clip}\left(\frac{\text{EMA of gradients}}{\text{EMA of diag Hessian estimate}}, \rho\right)$$
-

| Optim            | 類型                             | 額外狀態      | 特色                 | 何時試                                                     |
|------------------|--------------------------------|-----------|--------------------|---------------------------------------------------------|
| <b>AdamW</b>     | 一階自適應 +<br><b>decoupled wd</b> | m, v      | 預設穩、調參直覺           | Transformer/LL<br>M/ViT 基線<br>( <a href="#">arXiv</a> ) |
| <b>Adafactor</b> | 省記憶體（行列分<br>解二階）               | 低         | 近 Adam 效果、顯<br>存省  | 大模型/顯存緊<br>(T5 等) ( <a href="#">arXiv</a> )             |
| <b>Lion</b>      | 符號+動量（無 v）                     | m         | 省記憶體；大<br>batch 友好 | ViT/對比學習/<br>微調可試 ( <a href="#">arXiv</a> )             |
| <b>Sophia</b>    | 輕量二階（Hessian<br>對角估計+裁剪）       | m, diag-H | 步數減少報告<br>(LLM)    | 預訓練省步數<br>場景 ( <a href="#">arXiv</a> )                  |

# LEARNING RATE VS BATCH SIZE

# Learning rate vs Batch Size

- 增加batchsize為原來的N倍時，
  - 要保證經過同樣的樣本後更新的權重相等，按照線性縮放規則，學習率應該增加為原來的N倍[5]。
  - 但是如果要保證權重的方差不變，則學習率應該增加為原來的 $\sqrt{N}$ 倍[7]

Goyal P, Dollar P, Girshick R B, et al. Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour.[J]. arXiv: Computer Vision and Pattern Recognition, 2017.

Hoffer E, Hubara I, Soudry D. Train longer, generalize better: closing the generalization gap in large batch training of neural networks[C]//Advances in Neural Information Processing Systems. 2017: 1731-1741.