

LOW RANK APPROXIMATION

Low Rank Approximation

- Layer responses lie in a low-rank subspace
- Decompose a convolutional layer with d filters with filter size $k \times k \times c$ to
 - A layer with d' filters ($k \times k \times c$)
 - A layer with d filter ($1 \times 1 \times d'$)

$$A * B = C$$

$$B = U \Lambda V^T$$

$$C = A * (U \Lambda V^T) = (A * U) * \Lambda * V^T$$

SVD，CP分解、Tucker分解、Tensor Train分解、Block Term分解方法,現在很多網絡都是1x1的小的捲積，已經比較快了，用矩陣分解很難進行加速和壓縮




Figure 1. Illustration of the approximation. (a) An original layer with complexity $O(dk^2c)$. (b) An approximated layer with complexity reduced to $O(d'k^2c) + O(dd')$.

- Method
 - SVD, tucker, CP

speedup	rank sel.	Conv1	Conv2	Conv3	Conv4	Conv5	Conv6	Conv7	err. ↑ %
2×	no	32	110	199	219	219	219	219	1.18
2×	yes	32	83	182	211	239	237	253	0.93
2.4×	no	32	96	174	191	191	191	191	1.77
2.4×	yes	32	74	162	187	207	205	219	1.35
3×	no	32	77	139	153	153	153	153	2.56
3×	yes	32	62	138	149	166	162	167	2.34
4×	no	32	57	104	115	115	115	115	4.32
4×	yes	32	50	112	114	122	117	119	4.20
5×	no	32	46	83	92	92	92	92	6.53
5×	yes	32	41	94	93	98	92	90	6.47

低秩估計方法存在一個待解決的問題，就是保留多少秩是不明確的

WINOGRAD TRANSFORMATION

Convolution

Definition: Convolution in the time domain is equivalent to pointwise multiply in the frequency domain.

$$f * g = \mathcal{F}^{-1} \{ \mathcal{F}\{f\} \cdot \mathcal{F}\{g\} \}$$

$\mathcal{F}\{f\}$ and $\mathcal{F}\{g\}$ are the Fourier transforms of f and g
The asterisk denotes convolution, not multiplication.

$$f = [0 \ 0 \ 1 \ 2]$$

$$g = [10 \ 20 \ 30]$$

$$f * g = [30 \ 80]$$

Dot Product Approach

$$out[0] = 0 * 10 + 0 * 20 + 1 * 20 = 30$$

$$out[1] = 0 * 10 + 1 * 20 + 2 * 20 = 80$$

- 6 Multiplications
- 4 Additions

Shmuel Winograd (Winograd FFTs)



$$F(2, 3) = \begin{bmatrix} d_0 & d_1 & d_2 \\ d_1 & d_2 & d_3 \end{bmatrix} \begin{bmatrix} g_0 \\ g_1 \\ g_2 \end{bmatrix} = \begin{bmatrix} m_1 + m_2 + m_3 \\ m_2 - m_3 - m_4 \end{bmatrix} \quad (5)$$

where

$$\begin{aligned} m_1 &= (d_0 - d_2)g_0 & m_2 &= (d_1 + d_2)\frac{g_0 + g_1 + g_2}{2} \\ m_4 &= (d_1 - d_3)g_2 & m_3 &= (d_2 - d_1)\frac{g_0 - g_1 + g_2}{2} \end{aligned}$$

- Data (d 's): 4 ADDs
 - Filter (g 's): 3 ADDs, 2 MULs
 - Outputs (m 's): **4 MULs, 4 ADDs**

Winograd Convolution

- Winograd's minimal filtering algorithm for small filter sizes (3x3)
 - computing m outputs with an r -tap FIR filter, which we call $F(m, r)$, requires $\mu(F(m, r)) = m + r - 1$ multiplications
 - minimal 2D algorithms: for computing $m \times n$ outputs with an $r \times s$ filter, which we call $F(m \times n, r \times s)$. These require $\mu(F(m \times n, r \times s)) = \mu(F(m, r))\mu(F(n, s)) = (m + r - 1)(n + s - 1)$


$$\mu(F(2, 3)) = 2 + 3 - 1 = 4 \quad 4 \text{ mul v.s. } 6 \text{ mul}$$

$F(2 \times 2, 3 \times 3)$ uses $4 \times 4 = 16$ multiplications
 standard algorithm uses $2 \times 2 \times 3 \times 3 = 36$
 an arithmetic complexity reduction of $36/16 = 2.25$

N	cuDNN		$F(2 \times 2, 3 \times 3)$		Speedup
	msec	TFLOPS	msec	TFLOPS	
1	12.52	3.12	5.55	7.03	2.26X
2	20.36	3.83	9.89	7.89	2.06X
4	104.70	1.49	17.72	8.81	5.91X
8	241.21	1.29	33.11	9.43	7.28X
16	203.09	3.07	65.79	9.49	3.09X
32	237.05	5.27	132.36	9.43	1.79X
64	394.05	6.34	266.48	9.37	1.48X

Table 5. cuDNN versus $F(2 \times 2, 3 \times 3)$ performance on VGG Network E with fp32 data. Throughput is measured in Effective TFLOPS, the ratio of direct algorithm GFLOPs to run time.

Winograd Convolution



Producing 4 output pixels:


Direct Convolution:

- $4 * 9 = 36$ multiplications (**1x**)

Winograd convolution:

- $4 * 4 = 16$ multiplications (**2.25x less**)

Winograd Convolution v.s. Sparse



Producing 4 output pixels:


Direct Convolution:

- $4 \times 9 = 36$ multiplications (**1x**)
- sparse weight [NIPS'15] (**3x**)
- sparse activation (relu) (**3x**)
- Overall saving: **9x**

Winograd convolution:

- $4 \times 4 = 16$ multiplications (**2.25x less**)
- dense weight (**1x**)
- dense activation (**1x**)
- Overall saving: **2.25x**

Winograd Convolution + Sparse



Producing 4 output pixels:

Direct Convolution:

- $4 \times 9 = 36$ multiplications (**1x**)
- sparse weight [NIPS'15] (**3x**)
- sparse activation (relu) (**3x**)
- Overall saving: **9x**

Winograd convolution:

- $4 \times 4 = 16$ multiplications (**2.25x less**)
- sparse weight (**2.5x**)
- dense activation (**2.25x**)
- Overall saving: **12x**

- Pruned Winograd-transformed weights
- Moving ReLU to the Winograd domain






Figure 1: Combining Winograd convolution with sparse weights and activations. (a) Original Winograd-based convolution proposed by Lavin (2015) fills in the non-zeros in both the weights and activations. (b) Pruning the 4×4 transformed kernel restores sparsity to the weights. (c) Moving the ReLU layer after Winograd transformation also restores sparsity to the activations.²⁴

Winograd Convolution + Sparse



overall activation density of 41:1% compared to 36:9% density for the spatial activations

Compressed neural architecture utilizing dimensionality reduction and quantization

